



dispel4py: An Open-Source Python library for Data-Intensive Seismology

Rosa Filgueira, Amrey Krause, Alessandro Spinuso,
Iraklis Klampanos, Peter Danecek, and Malcolm Atkinson

Roadmap

- + Introduction
- + Dispel4py features
- + dispel4py basic concepts
- + dispel4py advance concepts
- + Verce project
- + Seismic ambient noise cross-correlation
- + Conclusions and future work

Introduction – What it is dispel4py ?

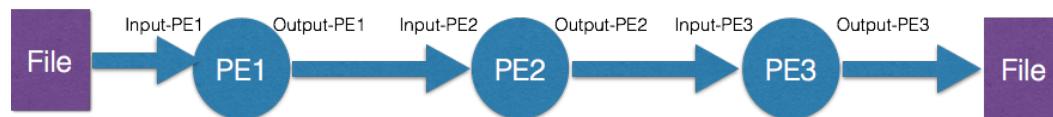
- + New user-friendly tool
- + Develop scientific methods and applications on local machines
- + Run them at scale on a wide range of computing resources without making changes

dispel4py features

- + **Stream-based**
 - + Tasks are connected by streams and not by intermediate files
 - + Multiple streams in & out
 - + Optimization based on avoiding IO
- + **Maps workflows dynamically** onto multiple enactment systems
- + **Python** language for describing tasks and connections
- + Reproducible and open scientific methods via a **registry and provenance mechanisms.**

dispel4py basic concepts – Processing element

- + PEs represent the basic computational blocks of any dispel4Py workflow: algorithm, service, data transformation
- + Shared: Storing them into the registry
- + PEs can be seen as “Lego bricks” of tasks. Users can assemble them into workflow as they wish.
- + General PE features
 - + Consumes any number and types of input streams
 - + Produce any number and types of output streams
 - + It does some processing → computational activity



dispel4py basic concepts – Instance and graph

+ Graph

- + Topology of the workflow:
connections among PEs

+ Instance

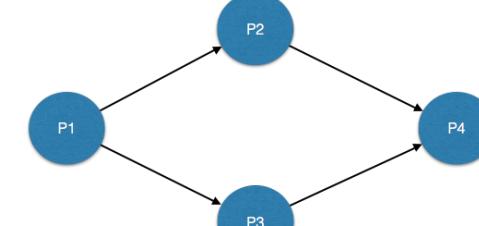
- + Executable copy of a PE that
runs in a process.
- + Each PE is translated into one
or more instances in run-time

+ Example of graphs

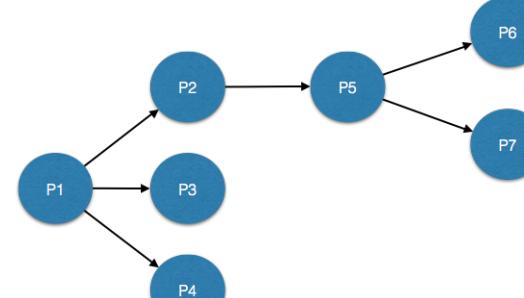
Pipeline



Split & Merge



Tree



dispel4py basic concepts – Composite PE and partition

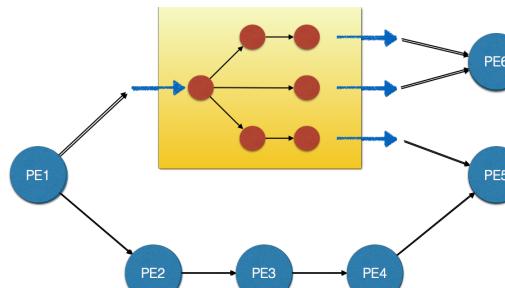
+ Composite PE

- + Sub-workflow in a PE
- + Hides the complexity of an underlying process
- + Treated like any other PE

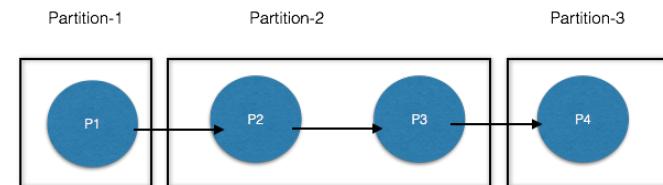
+ Partition

- + PEs wrapped together
- + Run several PEs in a single process

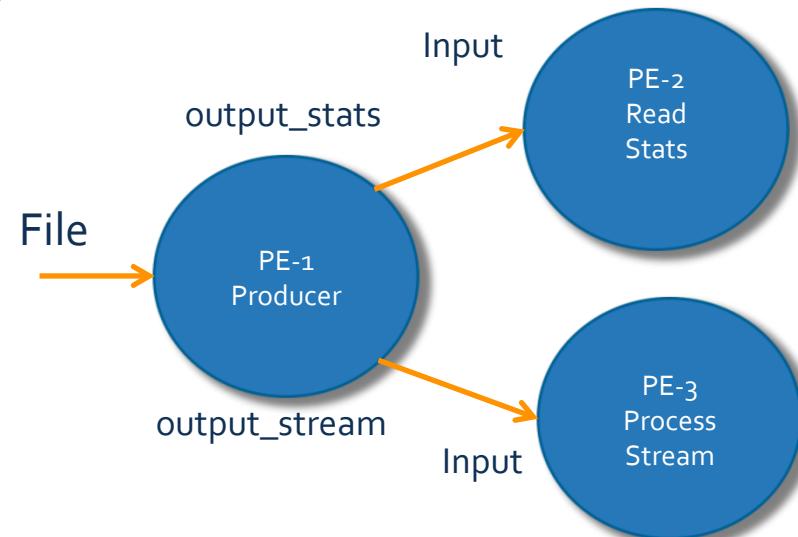
+ Example of Composite PE



+ Example of Partition



dispel4py basic concepts – Graph example



Users only have to implement:
PEs
Connections

```
from dispel4py.workflow_graph import WorkflowGraph
pe1 = Producer()
pe2 = ReadStats()
pe3 = ProcessStream()
graph= WorkflowGraph()
graph.connect (pe1, 'output_stats', pe2, 'input')
graph.connect (pe1, 'output_stream', pe3, 'input')
```

dispel4py basic concepts --

Example of a PE

```
from dispel4py.core import GenericPE  
from obspy import read
```

Class Producer(GenericPE):

""

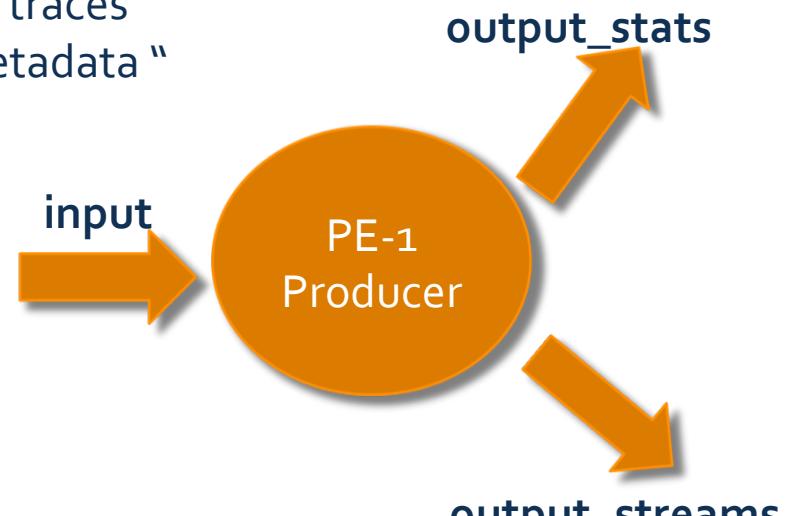
This PE reads a file that contains seismological traces
and returns two outputs: obspy stream and metadata "

def __init__(self):

```
    GenericPE.__init__(self)  
    self._add_input('input')  
    self._add_output('output_stream')  
    self._add_output('output_stats')
```

def process(self, inputs):

```
    data = inputs['input']  
    st = read(data)  
    return {'output_stream': st, 'output_stats': st[0].stats}
```



dispel4py – Advance concepts

- + Mappings
 - + Automatic and efficient parallelization to different parallel engines
 - + Without any cost for users
- + Provenance
 - + Runtime analysis of the provenance information collected
 - + Interactive validation of results
- + Registry
 - + Store workflows (and their PEs) with annotations

dispel4py advance concepts – Mappings

- + **Sequential**
 - + Sequential mapping for local testing
 - + Ideal for local resources: Laptops and Desktops
- + **Multiprocessing**
 - + Python's multiprocessing library
 - + Ideal for shared memory resources
- + **MPI**
 - + Distributed Memory, message-passing parallel programming model
 - + Ideal for HPC clusters
- + **STORM**
 - + Distributed Real-Time computation System
 - + Fault-tolerant and scalable

Run dispel4py workflows

Sequential mapping

```
>> dispel4py simple <name_dispy_graph>
```

Multi-process mapping

```
>> dispel4py multi <name_dispy_graph> -n <number mpi_processes>
```

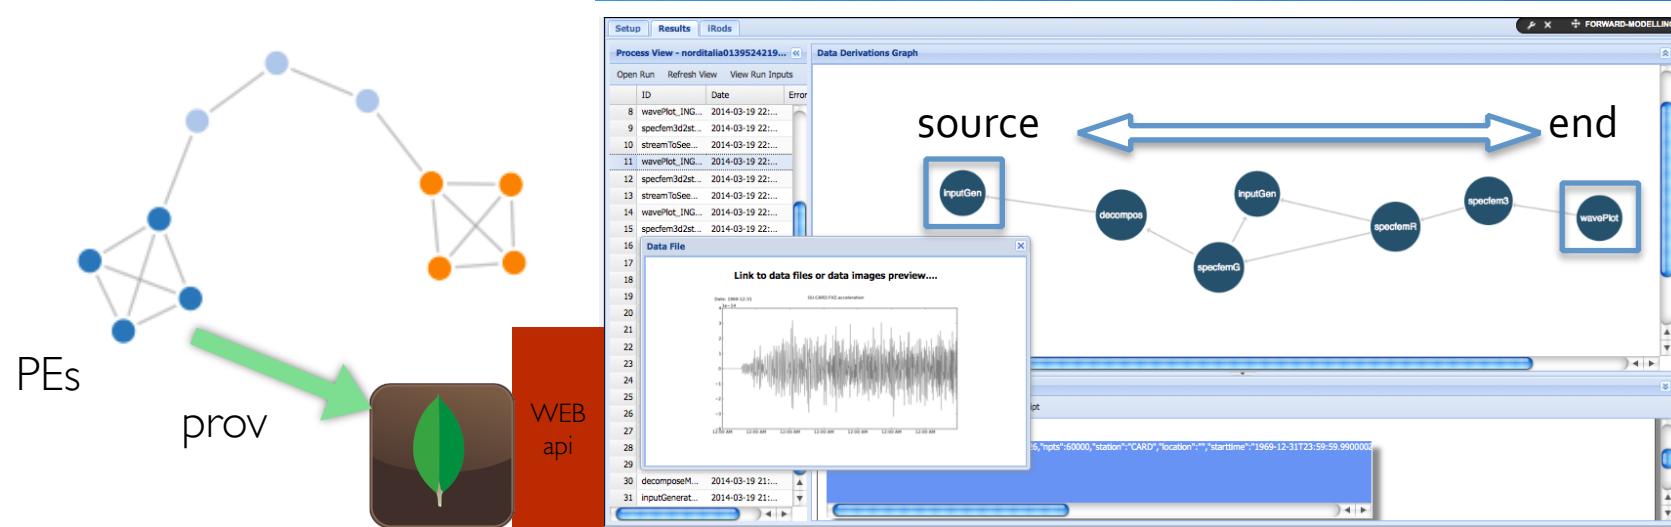
MPI mapping

```
>> mpiexec -n <number mpi_processes> dispel4py mpi <name_dispy_graph>
```

STORM mapping

```
>> python storm_submission.py <name_dispy_graph> <-m mode:  
[remote, local, create]>
```

dispel4py advance concepts – Provenance



Users can select which metadata to store
 Searches over products metadata within and across runs
 Data download and preview
 Capturing of Errors for Diagnostic purposes
Data Fabric: Multi directional navigations across data dependencies
 W3C PROV-DM as reference model.

dispel4py advance concepts – Registry

The screenshot shows a web-based administrative interface for managing Process Elements (PEs) in a dispel4py workspace. The URL in the browser is 127.0.0.1.

VERCE Registry Admin

Welcome, admin. Change password / Log out

Home > Vercereg > PEs > Add PE

Add PE

Workspace: admin: SampleWorkspace

Pckg: test.package

Name: SamplePE

User: admin

Creation date: Date: 2015-03-23 | Today | Time: 09:51:04 | Now |

Note: You are 2 hours ahead of server time.

Description: A sample dispel4py PE definition.
PE definitions can potentially have multiple implementations, which can be provided at a later time.

Kind: Abstract

Connections

Kind	Name	S type	D type	Comment	Is array	Modifiers	Delete?
In	Input1	str			<input checked="" type="checkbox"/>	modifier1:modifier2	<input type="button" value="Delete"/>
Out	Output1	int			<input type="checkbox"/>	modifier3	<input type="button" value="Delete"/>

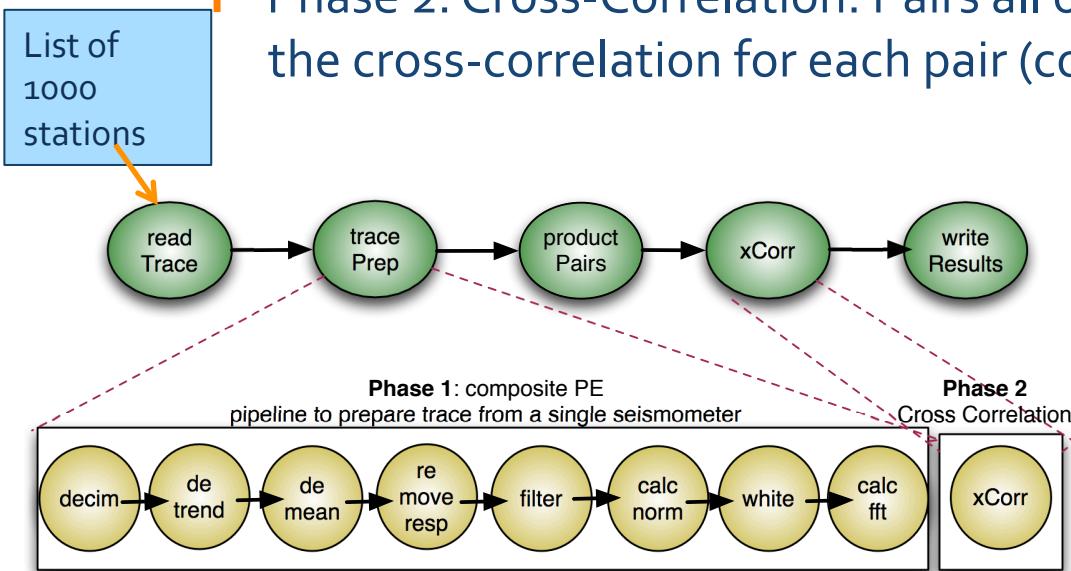
[+ Add another Connection](#)

VERCE project

- + The VERCE project provide a framework to the seismological community to exploit the increasingly large volume of seismological data :
 - + Support data-intensive and HPC applications
 - + e-Science Gateway for submitting applications
 - + Distributed and diversified data sources
 - + Distributed HPC resources on Grid, Cloud and HPC clusters
- + Use cases – dispel4py :
 - + **Seismic Noise Cross-Correlation**
 - + MISFIT

Seismic ambient noise cross-correlation – 1000 stations

- + Data intensive problem and it is commonly used in seismology
 - + Phase 1- Preprocess: Time series data (traces) from seismic stations are preprocessed in parallel
 - + Phase 2: Cross-Correlation: Pairs all of the stations and calculates the cross-correlation for each pair (complexity $O(n^2)$).



Input data:

1000 stations as input data (150 MB)

Output data:

499,500 cross-correlations (39GB)

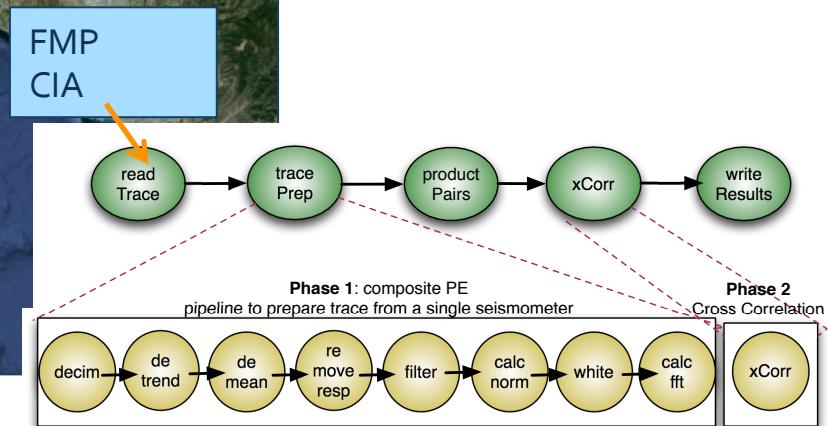
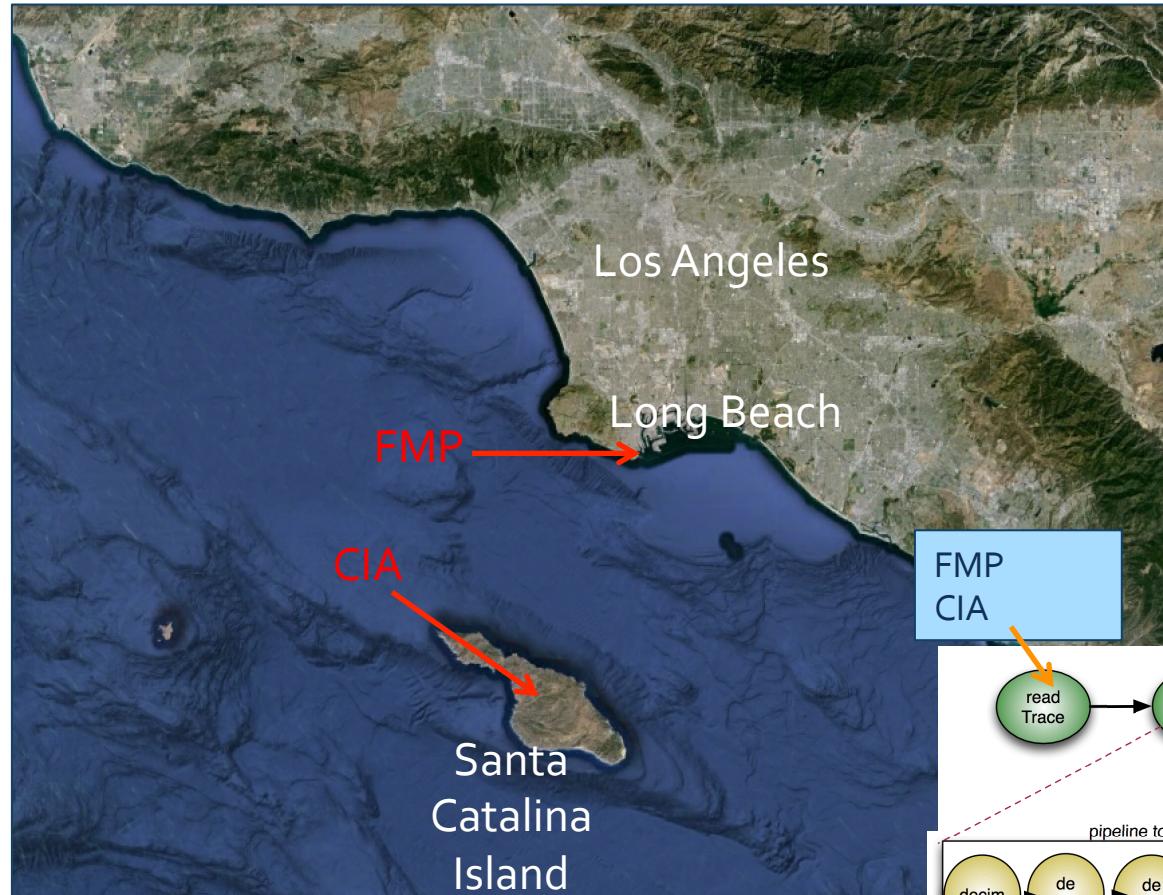
Evaluations – Computing resources

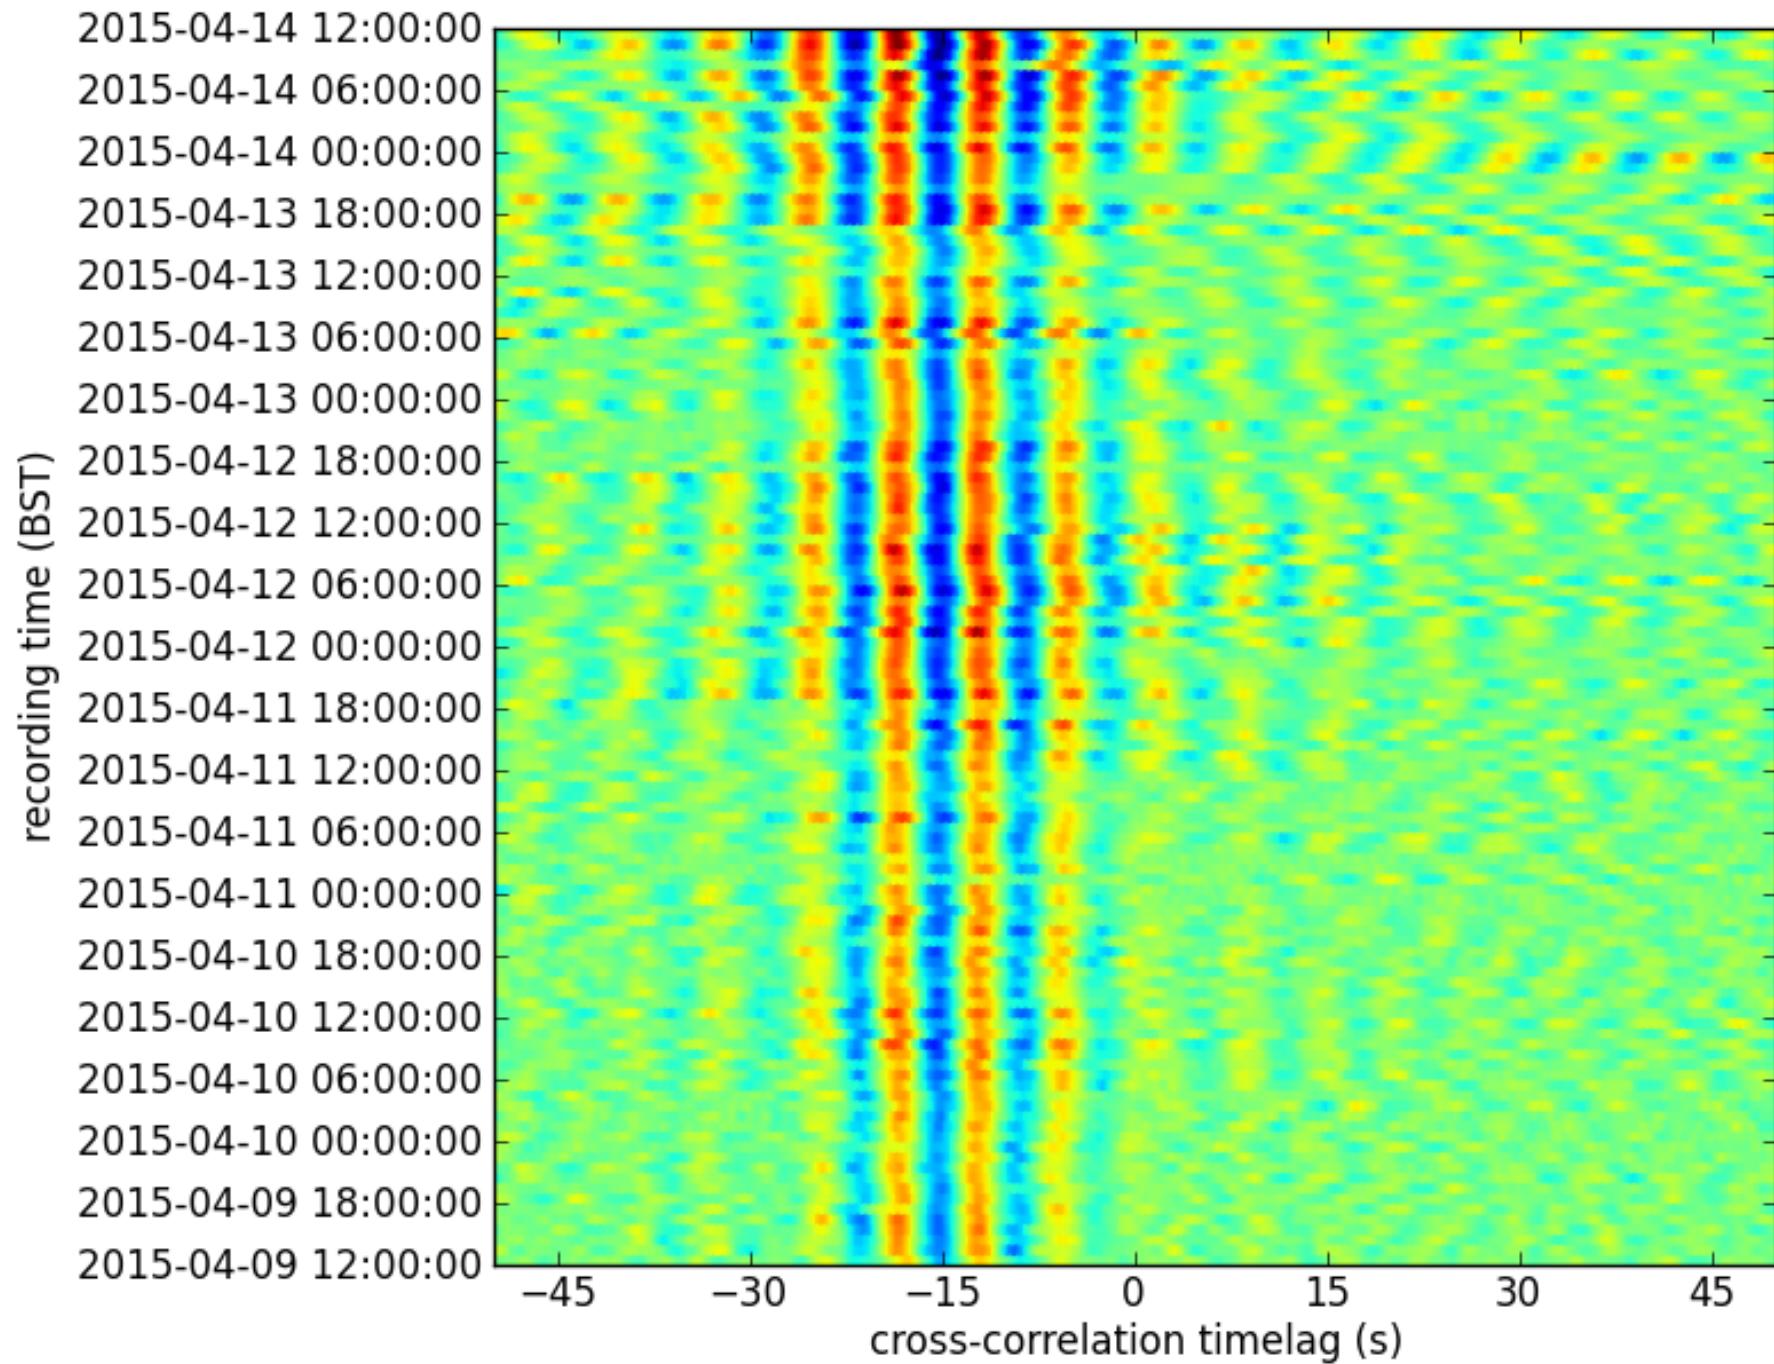
Computing Resources	Terracorrelator	SuperMuc	Amazon EC2	EDIM1
Type	Shared-memory	Cluster	Cloud	Cloud
Enactment Systems	MPI, multi	MPI, multi	MPI, Storm, multi	MPI, Storm, multi
Nodes	1	16	18	14
Cores per Node	32	16	2	4
Total Cores	32	256	36	14
Memory	2TB	32GB	4GB	3GB

Evaluations – Measures in seconds

Computing Resources / Enactment System	Terracorrelator (32 cores)	SuperMuc (256 cores)	Amazon (36 cores)	EDIM1 (14 cores)
MPI	2510.22 (~42minutes)	1093.16 (~19minutes)	16862.73 (~5hours)	38656.94 (~11 hours)
multi	2665.10 (~44minutes)		27898.89 (~8 hours)	
Storm				120077.123 (~33 hours)

Near real-time continuous correlations of two stations:
CI-FMP and CI-CIA
from the Caltech Regional Seismic Network





Conclusions and future Work

- + Novel Python library for **streaming and data-intensive processing**
- + **Novelty**
 - + Users express their computational activities
 - + Same workflow can be automatically executed in several parallel systems
 - + dispel4py is easy to use, allows users to share , re-use Pes and collect provenance information.
- + **Future**
 - + Support for PE failures
 - + New mapping: Apache Spark
 - + diagnostic tools to select the best computing resource and mapping

Thanks & Questions

- + Documentation
 - + <http://dispel4py.org>
- + Source code
 - + <https://github.com/dispel4py/dispel4py>
- + Installation
 - + **pip install dispel4py**
- + Contact emails
 - + Rosa Filgueira: rosa.filgueira@ed.ac.uk
 - + Amy Krause: akrause@staffmail.ed.ac.uk
 - + Alessandro Spinuso: spinuso@knmi.nl