

Virtual Earthquake and seismology Research Community e-science environment in Europe
Project 283543 – FP7-INFRASTRUCTURES-2011-2 – www.verce.eu – info@verce.eu

dispel4py: An Agile Framework for Data-Intensive eScience

11Th IEEE eScience 2015
2nd September 2015



Roadmap

Introduction

dispel4py features

dispel4py basic concepts

dispel4py advance concepts

dispel4py workflows

Evaluations

Current work

Conclusions and future work

What it is dispel4py ?

New user-friendly tool

Develop scientific methods and applications on local machines

Run them at scale on a wide range of computing resources without making changes

What it is dispel4py ?

open source project: www.dispel4py.org & <https://github.com/dispel4py/dispel4py> & pip install dispel4py

publications:

DISCS-2014 workshop, SC14, 2014

IJHPCA journal, "Data-Intensive High Performance Computing" Special Issue, 2015

11th IEEE eScience Conference, 2015

Book Chapter in "Conquering Big Data Using High Performance", 2015

any users:

Computational Seismologists, VERCE project.

Astrophysicists

trainings:

General audience, OSDC PIRE, 2015, UvA

VERCE trainings: Munich and Liverpool, 2015

any contributors:

University of Edinburgh: EPCC, DIR Group, Master Students

KNMI

dispel4py features

Stream-based

Tasks are connected by streams. Not by intermediate files

Optimization based on avoiding IO.

Multiple streams in & out

Data-flow oriented

Data is transformed among the tasks

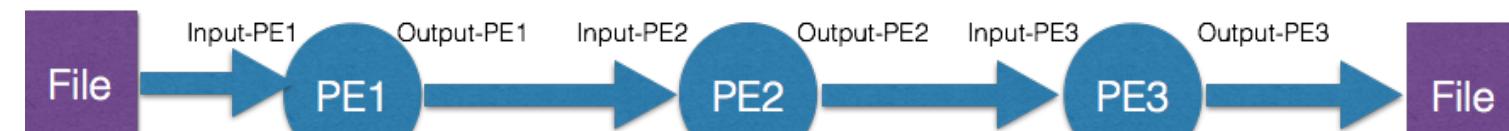
Maps workflows dynamically onto multiple enactment systems

Python language for describing tasks and connections

Encourage the reproducibility & re-usability of workflows:

Processing element

- PEs represent the basic computational:
Algorithm, service, data transformation
- Shared: Storing them into the registry
- PEs ~ “Lego bricks” of tasks.
Users can assemble them into workflow as they wish.
- Consume/Produce any number
and types of input/output streams



Instance and graph

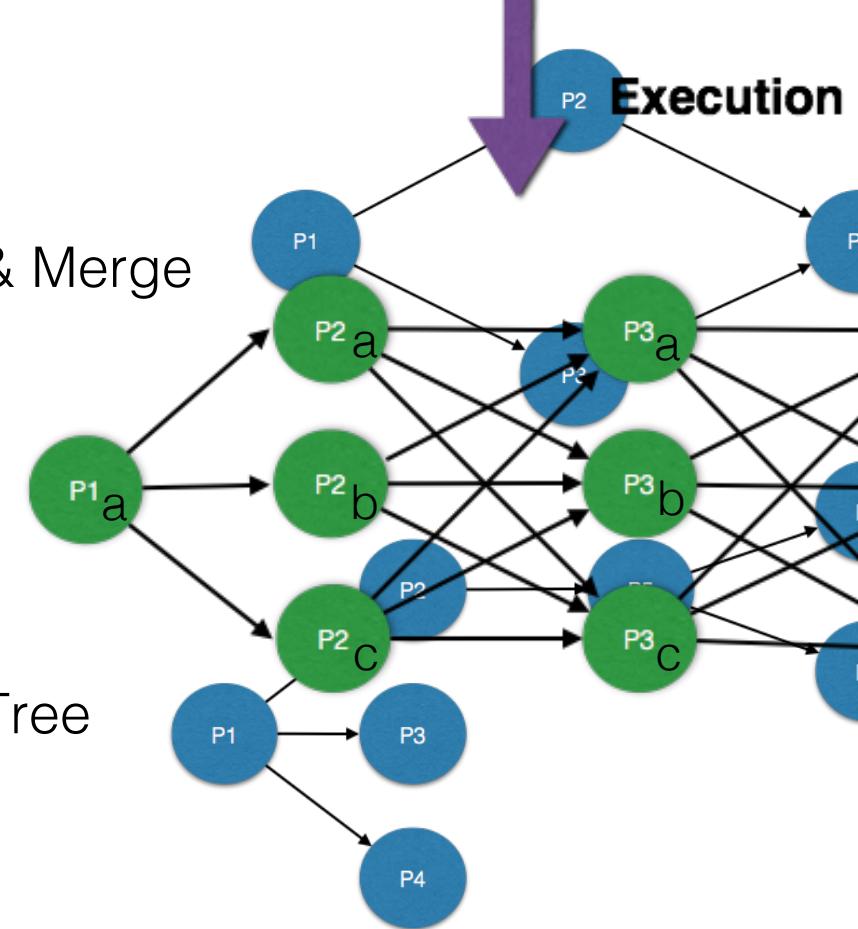
- Graph
 - Topology of the workflow: connections among PEs
 - It is what user need to “think” about.
 - Abstract workflow
- Instance
 - Executable copy of a PE that runs in a process.
 - Each PE is translated into one or more instances in run-time
 - dispel4py does it for you
 - Concrete workflow

+ Example of graphs

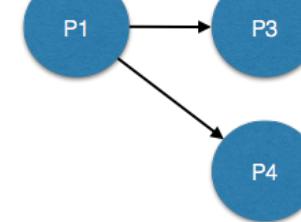
Pipeline



Split & Merge



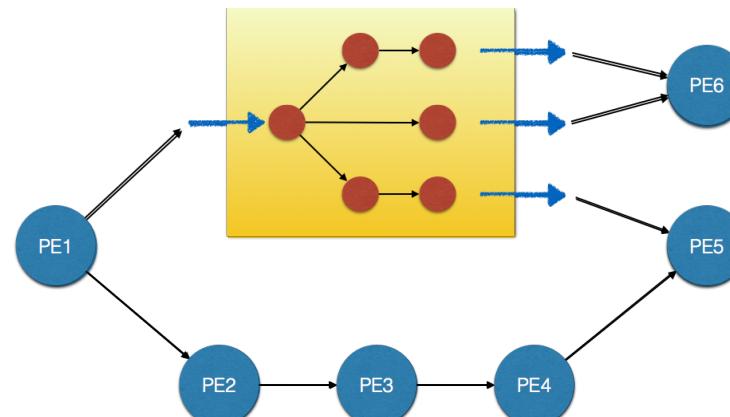
Tree



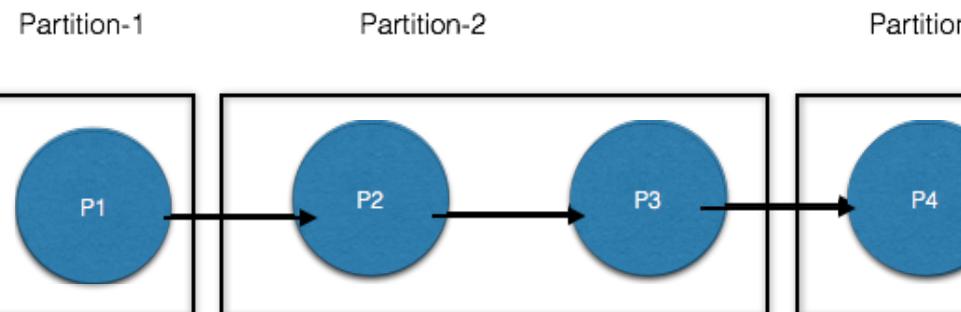
Composite PE and partition

- Composite PE
 - Sub-workflow in a PE
 - Hides the complexity
 - Treated like any other PE
- Partition
 - PEs wrapped together
 - Run several PEs in a 1 process

+ Example of Composite PE



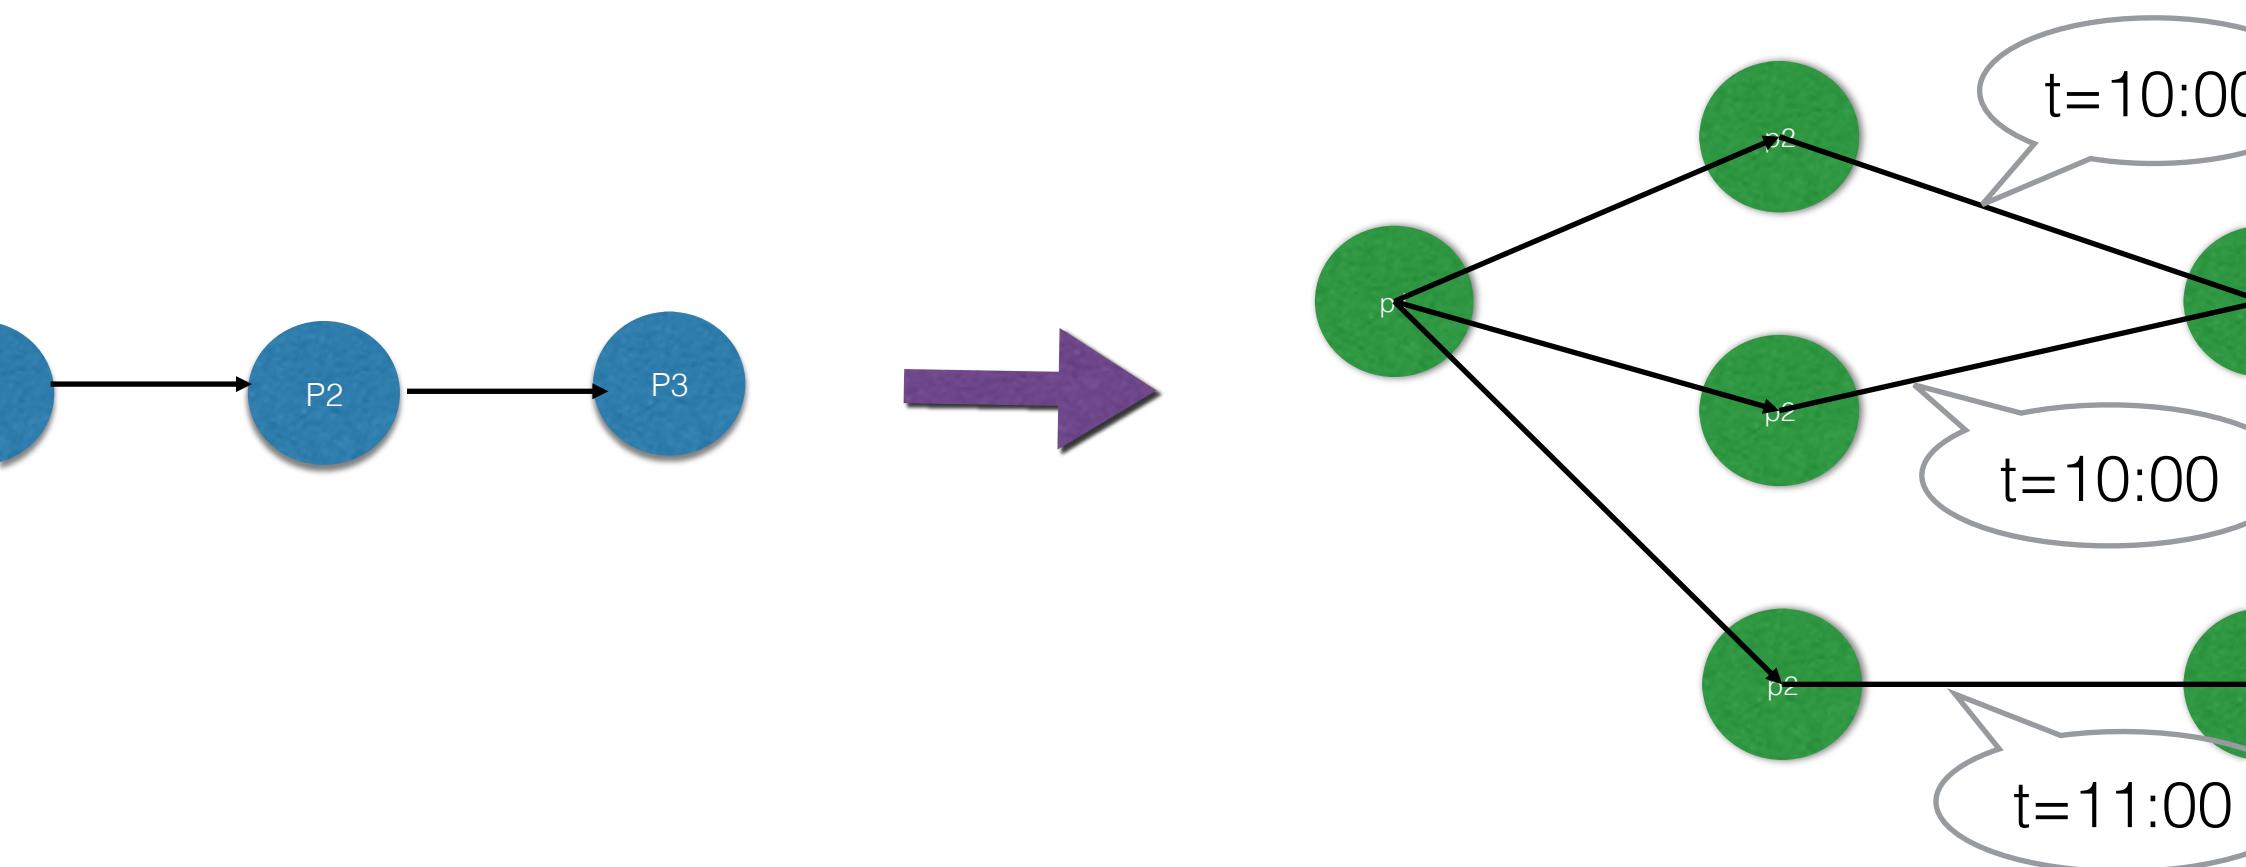
+ Example of Partition



Groupings

Grouping by” a feature (MapReduce)

data items that satisfy the same feature are guaranteed to be delivered to the same **instance** of a PE



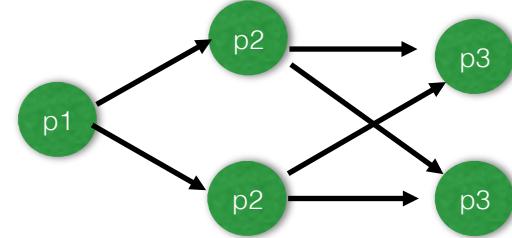
Groupings

-To-All



P3 - grouping “all”:

Instances send copies of their output data to **all** the connected instances

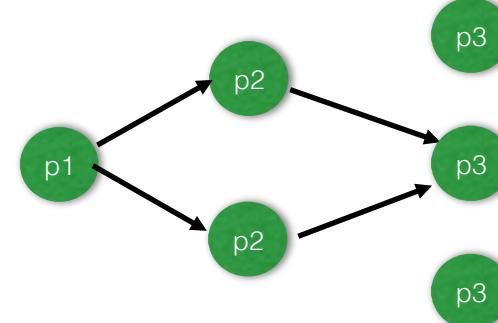


bal



P3 - grouping “global”:

All the instances send all the data to **one** connected instance



Example of a dispel4py workflow

```
from dispel4py.workflow_graph import WorkflowGraph
```

```
filterTweet ()  
counterHashTag ()  
counterLanguage ()  
statistics ()
```

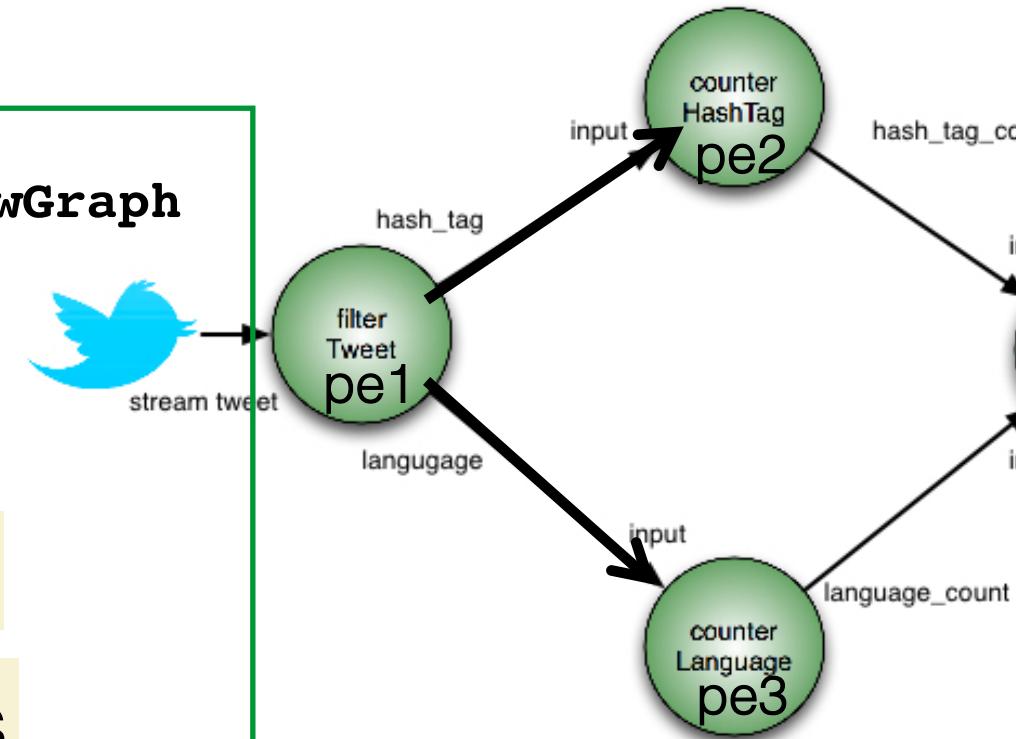
```
= WorkflowGraph ( )
```

```
connect(pe1,'hash_tag',pe2,'input') ←  
connect(pe1,'language',pe3,'input') ←  
connect(pe2,'hash_tag_count',pe4,'input1')  
connect(pe3,'language_count',pe4,'input2')
```

PEs objects

Graph

Connections



Users only have to implement:

- PEs
- Connections

Example of a concrete dispel4py workflow

1):

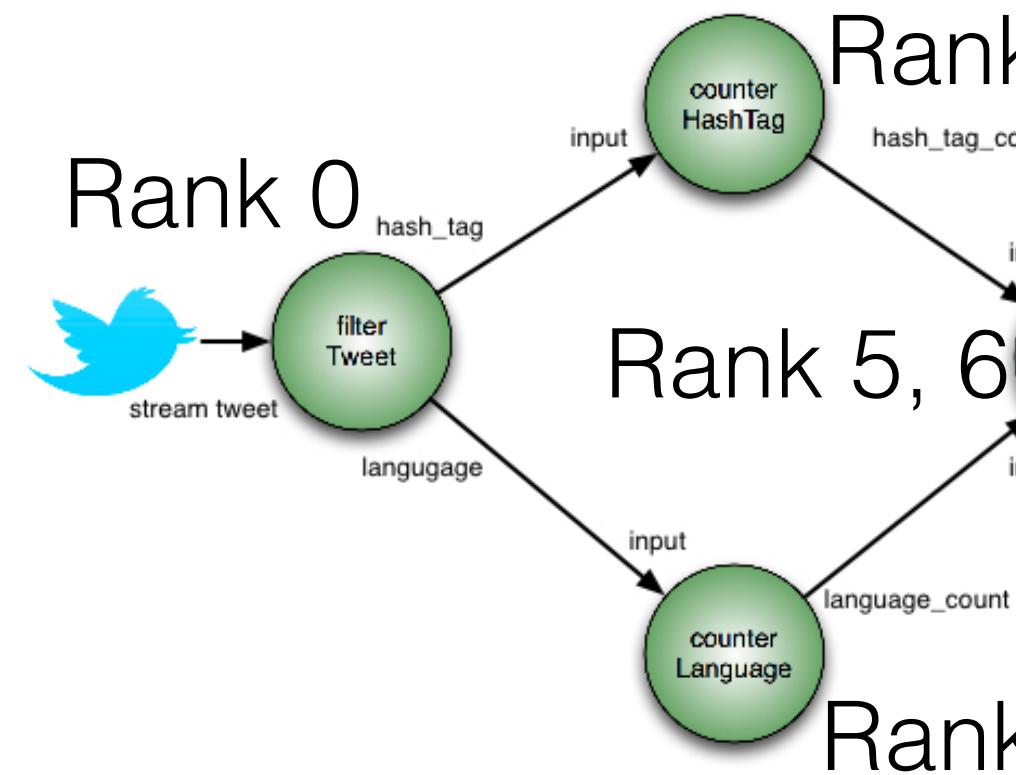
```
nk == 0:  
ad (data)  
sh_tag= Filter(data, hash_tag)  
nguage = Filter(data, language)  
I_Isend(hash_tag, rank → 1,2)  
I_Isend(language, rank → 3,4)
```

nk ==1 or rank == 2 :

```
I_Recv(input, 0->rank)  
sh_tag_count = Count(input)  
I_Isend(hash_tag_count, rank → 5,6)
```

nk == 3 or rank == 4:

```
nk == 5 or rank == 6:  
I_Recv(input1, 1,2 → rank)
```



Users only have to implement:

- PEs
- Connections

Example of a PE

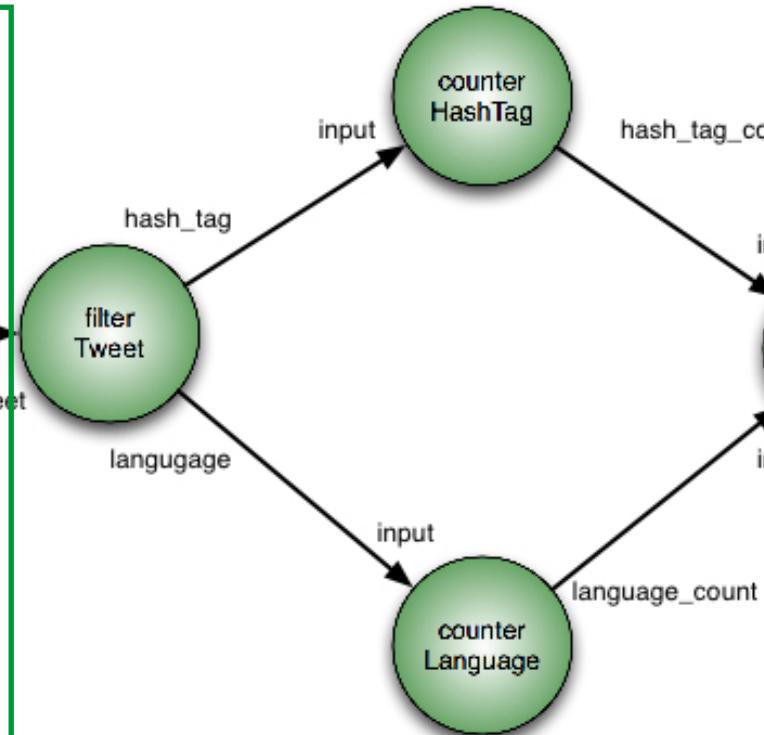
```
filterTweet(GenericPE):  
    def __init__(self):  
        GenericPE.init(self)  
        self.add_output('hash_tags')  
        self.add_output('language')
```

```
process(self, inputs):  
    twitterDataFile = inputs['input']  
    tweet_file = open(twitterDataFile)  
    for line in tweet_file:  
        tweet = json.loads(line)  
        language = ''  
        hashtags = []  
        language = tweet[u'lang'].encode('utf-8')  
        text = tweet[u'text'].encode('utf-8')  
        hashtags=re.findall(r"#(\w+)", text)  
  
        self.write('hash_tag', hashtags)  
        self.write('language', language)
```

Inputs & outputs



Logic
of PE



Users only have to implement:

- PEs
- Connections

Mappings

Sequential

- Sequential mapping for local testing
- Ideal for local resources: Laptops and Desktops

Multiprocessing

- Python's multiprocessing library
- Ideal for shared memory resources

MPI

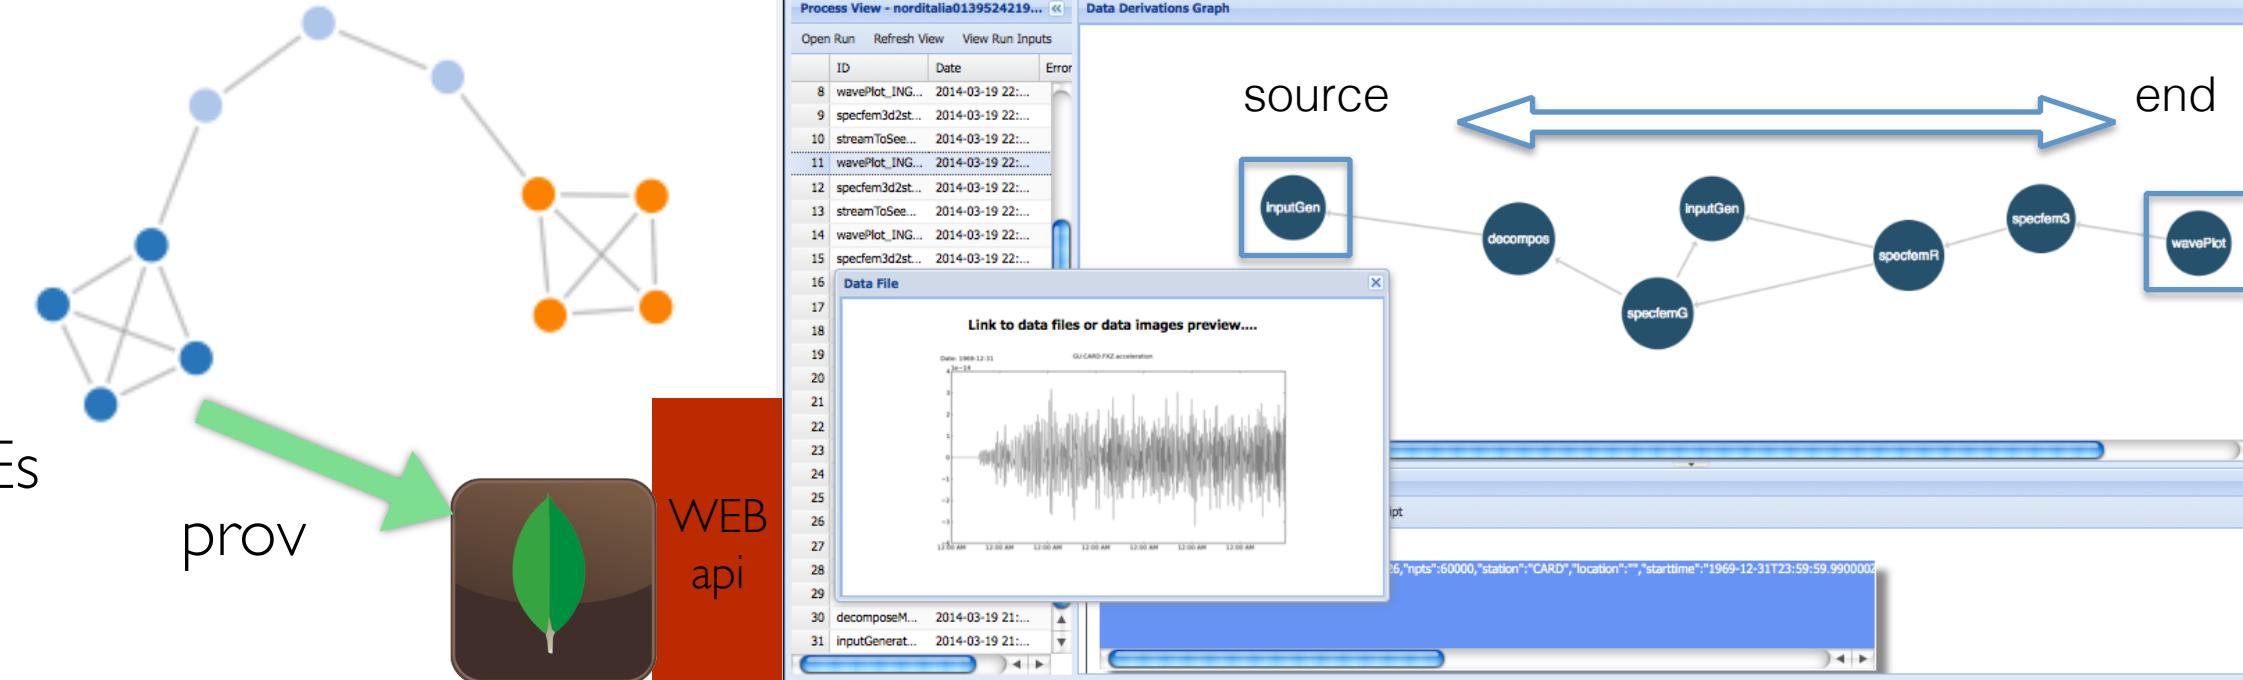
- Distributed Memory, message-passing parallel programming model
- Ideal for HPC clusters

STORM

- Distributed Real-Time computation System
- Fault-tolerant and scalable
- Runs all the time

Spark (Prototype)

Provenance



Users can select which metadata to store
Searches over products metadata within and across runs
Data download and preview
Capturing of Errors for Diagnostic purposes
Data Fabric: Multi directional navigations across data dependencies

Registry

127.0.0.1

Welcome, admin. Change password

Reg > PEs > Add PE

Name: admin: SampleWorkspace +

Implementation: test.package

Author: SamplePE

Owner: admin

Date: 2015-03-23 Today |

Time: 09:51:04 Now |

Note: You are 2 hours ahead of server time.

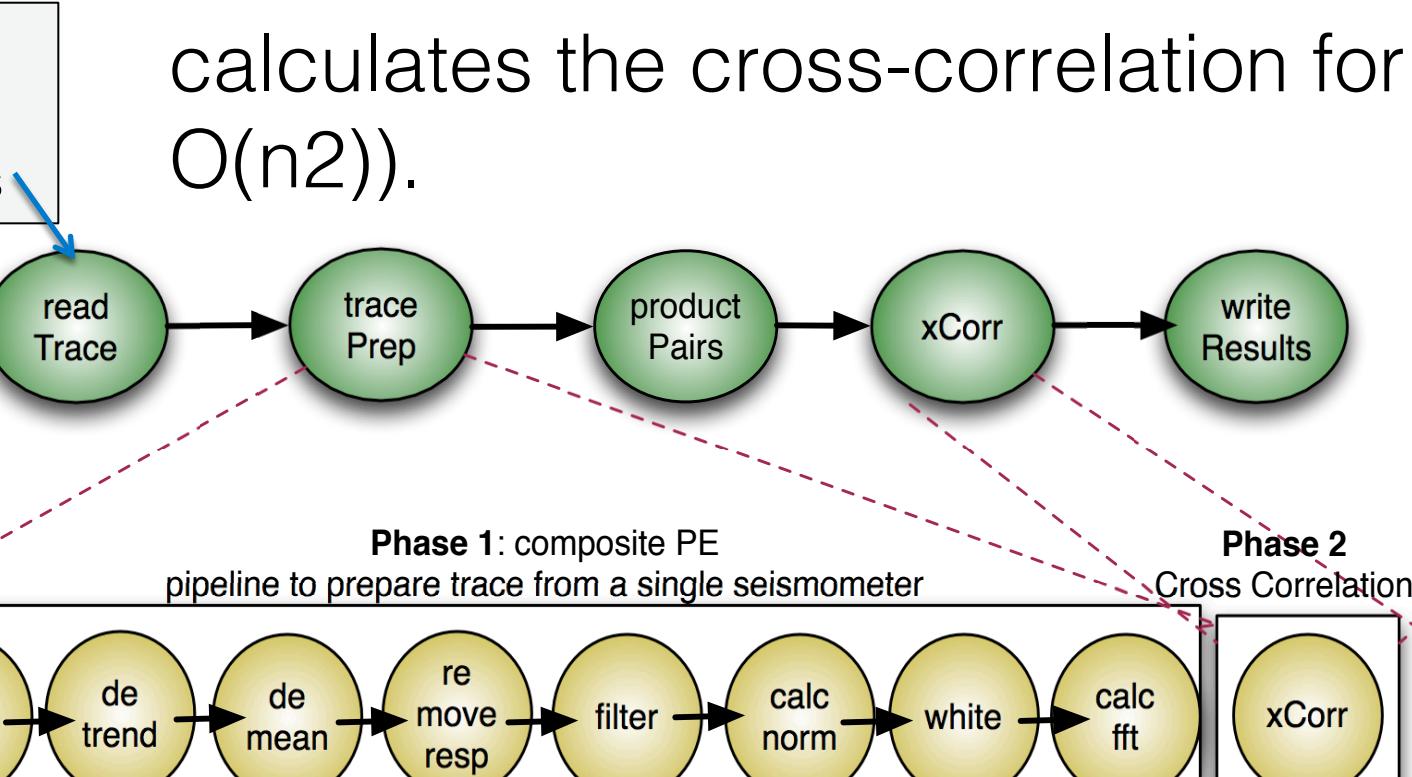
Description: A sample dispel4py PE definition.
PE definitions can potentially have multiple implementations, which can be provided at a later time.

Type: Abstract

Name	S type	D type	Comment	Is array	Modifiers
------	--------	--------	---------	----------	-----------

Seismology, Cross Correlation

- Data intensive problem and it is commonly used in seismology
- Phase 1- Preprocess: Time series data (traces) from seismic stations are preprocessed in parallel
- Phase 2: Cross-Correlation: Pairs all of the stations and calculates the cross-correlation for each pair (complexity $O(n^2)$).



Input data:

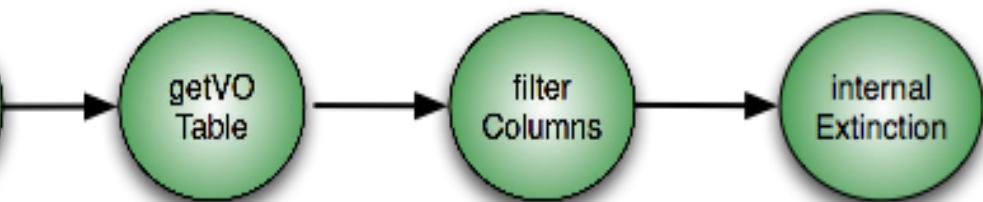
1000 stations as input data (1

Output data:

499,500 cross-correlations (39

Astrophysics, Internal Extinction

Calculates the Internal Extinction of the Galaxies from the AMIGA catalog.

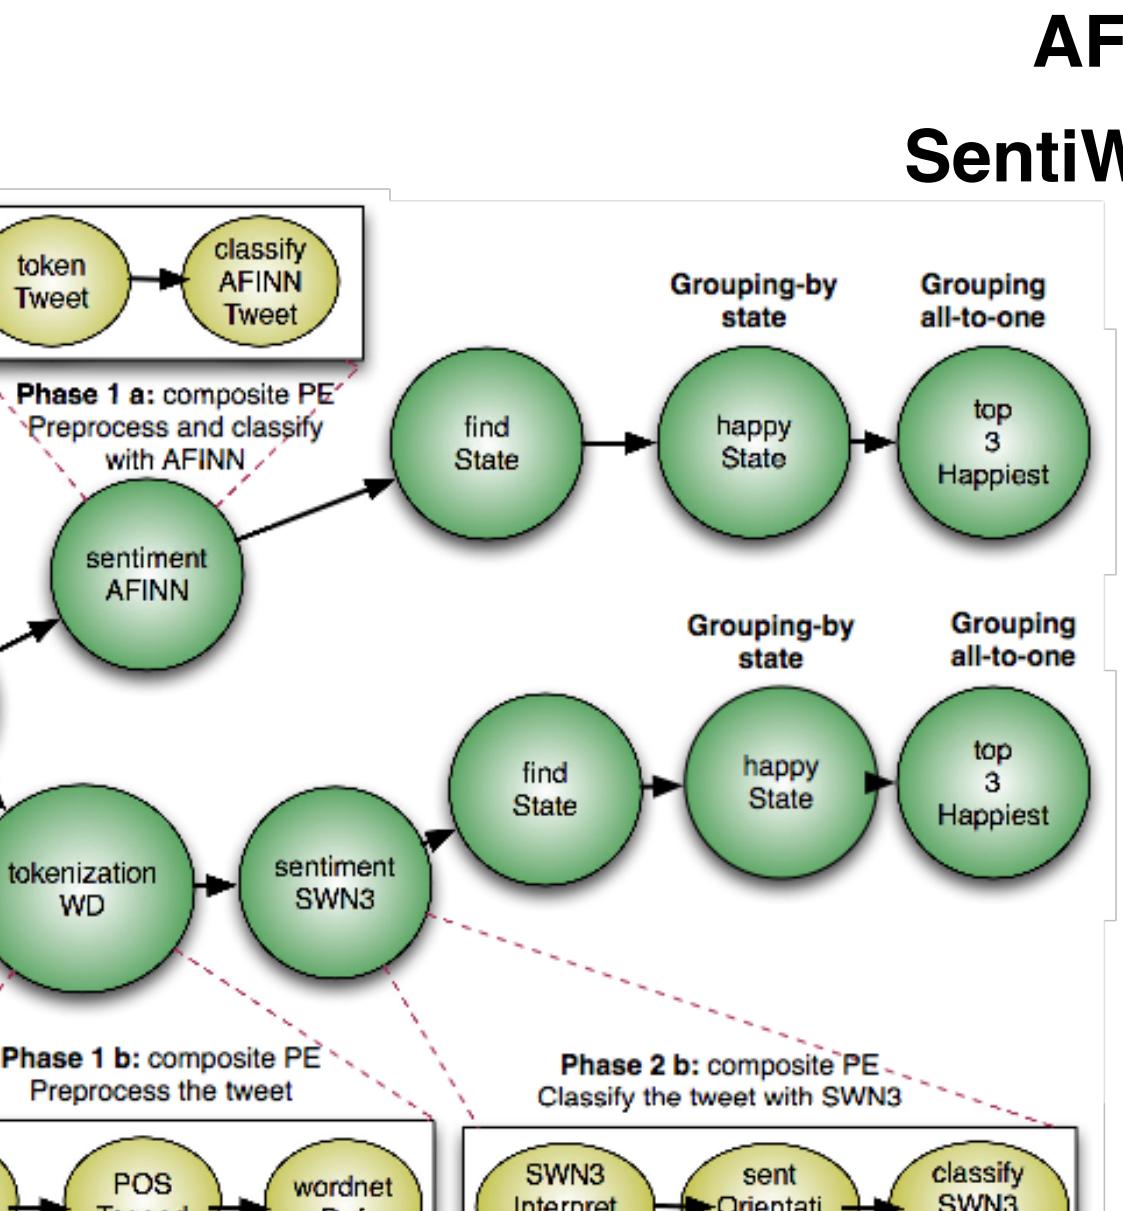


- The first PE **reads an input file (1051 galaxies)**:
 - **Right Ascension (Ra)**
 - **Declination (Dec)**
- The second PE **queries a VO service** for each galaxy
 - The results are filtered by filterColumns
 - Their internal extinction is calculated by internalExtinction PE.

_ext

Social Computing, Sentiment Analysis

basic sentiment analyses (126,826 tweets (500MB)) by applying two lexicons:



AFINN

SentiWordNet

- The `findState` PE searches the US state
- The `HappyState` PE aggregates the sentiment scores of tweets from the same state
- The last PE determines which are the top three happiest states.

sentiment

Computing resources

Computing sources	Terracorrelator	SuperMuc	Amazon EC2	EDIM1
Type	Shared-memory	Cluster	Cloud	Cloud
Execution systems	MPI, multi	MPI, multi	MPI, Storm, multi	MPI, Storm, multi
Nodes	1	16	18	14
Cores per Node	32	16	2	4
Total Cores	32	256	36	14
Memory	2TB	32GB	4GB	3GB
Workflows	xcorr, int_ext,	xcorr,	xcorr	xcorr, int_ext

Performance measures

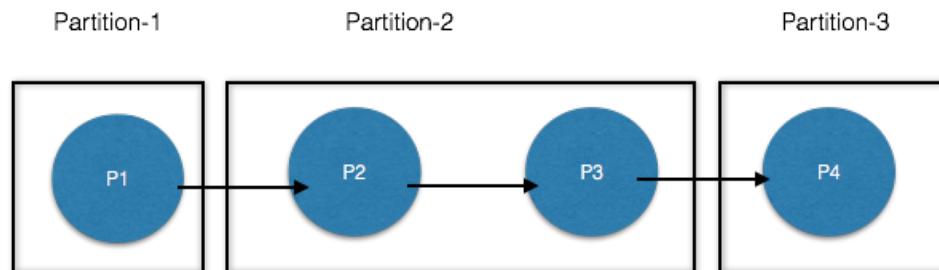
Terracorrelator (32 cores/Node)	SuperMuc (256 cores 16 cores/node)	Amazon (36 cores 2 cores/node)	EDIM1 (14 cores 4 cores/node)
1501.32 (~25minutes)	1093.16 (~19minutes)	16862.73 (~5hours)	38656.94 (~11 hours)
1332 .20 (~23minutes)			
		27898.89 (~8 hours)	120077.123 (~33 hours)
Terracorrelator (32 cores/node)	EDIM1 (14 cores, 4 cores/node)	Mode	Terracorrelator (32 cores/node)
31.60	96.12	MPI	925.04 (~15minutes)
14.50	101.2	multi	971.04 (~16minutes)
	30.2		

xcorr:
1000 stations
Input 150MB
Output 39GB

sentiment:
3954.63
(~1hour)

Run-time Stream Adaptive Compression Diagnosis tool

- Quality check and Failure detector
- How to partition the workflow automatically
- How many processes execute each partition



Conclusions and Future work

Python library for streaming and data-intensive processing

- Users express their computational activities
- Same workflow executed in several parallel systems
- Easy to use, open, and encourage sharing the methods & applications

Future work

- Support for PE failures
- Select the best computing resource and mapping

Installations and Links

- This is all you need:

pip install dispel4py

- Web site <http://dispel4py.org/>
- GitHub: <https://github.com/dispel4py/dispel4py>
- Documentation: <http://dispel4py.org/documentation/>

Thanks and Questions

Contact emails

- Rosa Filgueira: rosa.filgueira@ed.ac.uk
- Amy Krause: a.krause@epcc.ed.ac.uk
- Alessandro Spinuso: spinuso@knmi.nl
- Malcolm Atkinson: Malcolm.Atkinson@ed.ac.uk